

Disjoint-set Union

Eklavya Sharma

Contents

1	Problem	2
2	Forest algorithm	2
2.1	Performance with no optimizations	3
2.2	rank upper-bounds height	3
2.3	Performance when <code>union_by</code> is not <code>None</code>	4
2.4	Lower bound on time when <code>compress_path</code> is <code>False</code>	5
2.5	Both union-by-rank and path-compression	5
2.5.1	Alt-Ackermann function	5
2.5.2	level and iter	6
2.5.3	Potential function	7

All (pseudo-)code in this document is based on the python programming language.

1 Problem

In the Disjoint-set Union (DSU) problem, we are given a set S of n singleton sets, i.e. $S = \{\{i\} : 0 \leq i < n\}$.

We have to perform m operations on S . Each operation can modify S while maintaining these 2 invariants:

1. All elements of S are sets.
2. Every integer from 0 to $n-1$ lies in exactly one set in S .

Also, for every set $X \in S$, one of the elements of X will be known as the ‘representative of X ’, denoted as $\text{repr}(X)$.

Types of operations allowed

1. **find**(x): If $x \in X$, return $\text{repr}(X)$.
2. **union**(x, y): Let $x \in X$ and $y \in Y$. Then remove X and Y from S and add $X \cup Y$ to S .

union(x, y) is the only operation which can modify S . It is easy to see that **union**(x, y) maintains the 2 invariants.

2 Forest algorithm

The ‘forest algorithm’ for DSU maintains a forest $F = (V, E)$ of rooted trees where $V = \{i \in \mathbb{Z} : 0 \leq i < n\}$. Each tree in F corresponds to a set in S . The representative of a set is the root of the corresponding tree.

The forest is stored by keeping track of the parent of each vertex in an array **parent** of size n . If a vertex x has no parent, then **parent**[x] = x . The algorithm (optionally) maintains 2 additional arrays **rank** and **size**. **rank**[i] is an upper-bound on the height of vertex i and **size**[i] is the size of the subtree rooted at vertex i . Initially **parent**[i] = i , **rank**[i] = 0 and **size**[i] = 1 for all $0 \leq i < n$.

This algorithm offers 2 hyperparameters. These are optional optimizations for speeding up DSU.

1. **union_by**: can be **None**, **rank** or **size**.
2. **compress_path**: can be **False** or **True**.

This is how **find** and **union** are implemented:

```

1 def find(x):
2     if parent[x] == x:
3         return x
4     else:
5         r = find(parent[x])
6         if compress_path:
7             parent[x] = r
8         return r
9
10 def link(x, y):
11     parent[y] = x
12     size[x] += size[y]
13     rank[x] = max(rank[x], rank[y] + 1)
14
15 def union(x, y):
16     x = find(x)
17     y = find(y)
18     if union_by is not None and union_by[x] < union_by[y]:
19         x, y = y, x
20
21     link(x, y)
22     return x != y

```

2.1 Performance with no optimizations

Consider the following operations:

```

1 for i in range(1, n):
2     union(i, i-1)
3 for i in range(1, m - n):
4     find(0)

```

When `union_by` is `None`, `union(x, y)` makes the tree of y a subtree of x . Therefore, after all the `union` operations, the forest will be a single chain from 0 to $n-1$. If `compress_path` is `False`, each `find(0)` operation will take $\Theta(n)$ time. Each `union` operation takes $\Theta(1)$ time. Therefore, total time taken is $\Theta((m - n)n)$.

2.2 rank upper-bounds height

For a tree T , let $h(T)$ denote its height, $n(T)$ denote the number of nodes in it and $r(T) = \text{rank}(\text{repr}(T))$.

Theorem 1. $h(T) \leq r(T)$ throughout the algorithm.

Proof. Initially, $h(T) = r(T) = 0$ for every tree T .

In a `find` operation, the height of a tree can only reduce (it can reduce if `compress_path` is `True`, otherwise it doesn't change).

Suppose `link(x, y)` is called and $x \in X$ and $y \in Y$. Then Y is made a subtree of X . Let the resulting tree be Z . Suppose $h(X) \leq r(X)$ and $h(Y) \leq r(Y)$.

$$h(Z) = \max(h(X), h(Y) + 1) \leq \max(r(X), r(Y) + 1) = r(Z)$$

Since $h(T) \leq r(T)$ is initially true and remains true across `find` and `union` operations, $h(T) \leq r(T)$ is true for all trees across the entire DSU algorithm. \square

2.3 Performance when `union_by` is not `None`

Theorem 2. `union_by` \neq `None` $\implies \forall T, r(T) \leq \lg n(T)$.

Proof. Initially, $\forall T, r(T) = 0 = \lg 1 = \lg n(T)$.

`find` operations affect neither r nor n .

Suppose `link(x, y)` is called and $x \in X$ and $y \in Y$. Then Y is made a subtree of X . Let the resulting tree be Z . Suppose $r(X) \leq \lg n(X)$ and $r(Y) \leq \lg n(Y)$. $r(Z) = \max(r(X), 1 + r(Y))$ and $n(Z) = n(X) + n(Y)$.

Case 1: `union_by` = `size`

`union_by` = `size` $\implies n(Y) \leq n(X)$.

$$\begin{aligned} r(Z) &= \max(r(X), r(Y) + 1) \\ &\leq \max(\lg n(X), \lg n(Y) + 1) \\ &\leq \max(\lg n(X), \lg(2n(Y))) \\ &\leq \lg \max(n(X), 2n(Y)) \end{aligned}$$

$$n(X) \leq n(X) + n(Y) \text{ and } n(Y) \leq n(X) \Rightarrow 2n(Y) \leq n(X) + n(Y).$$

$$\implies r(Z) \leq \lg \max(n(X), 2n(Y)) \leq \lg(n(X) + n(Y)) = \lg n(Z)$$

Case 2: `union_by` = `rank`

`union_by` = `rank` $\implies r(Y) \leq r(X)$.

Case 2a: $r(Y) < r(X)$

$$\begin{aligned} r(Z) &= \max(r(X), 1 + r(Y)) = h(X) \\ &\leq \lg n(X) \leq \lg n(X) + n(Y) \leq \lg n(Z) \end{aligned}$$

Case 2b: $r(Y) = r(X)$

$$\begin{aligned} r(Z) &= \max(r(X), 1 + r(Y)) = 1 + r(Y) = 1 + r(X) \\ &\Rightarrow r(Z) \leq 1 + \lg n(Y) \wedge r(Z) \leq 1 + \lg n(X) \\ &\Rightarrow r(Z) \leq 1 + \min(\lg n(Y), \lg n(X)) \\ &\Rightarrow r(Z) \leq \lg(2 \min(n(X), n(Y))) \\ &\Rightarrow r(Z) \leq \lg(n(X) + n(Y)) = \lg n(Z) \end{aligned}$$

For both cases 1 and 2, $r(Z) \leq \lg n(Z)$. Therefore, `union` preserves the invariant $\forall T, r(T) \leq \lg n(T)$. \square

This means that any tree can have height at most $\lg n$. Therefore, **find** and **union** have a worst-case time complexity of $O(\lg n)$ and **link** has a worst-case time complexity of $O(1)$.

2.4 Lower bound on time when compress_path is False

When there is no path compression, we can lower bound the worst-case time complexity of **find**.

Consider these union operations:

```

1 for i in range(int(log2(n))):
2     for j in range(0, n, 1 << (i+1)):
3         union(j, j + (1 << i))

```

The body of the outer loop is called a round. There are $\lfloor \lg n \rfloor$ rounds.

Number of union operations:

$$\sum_{i=1}^{\lfloor \lg n \rfloor} \left\lfloor \frac{n}{2^i} \right\rfloor \leq n \sum_{i=1}^{\lfloor \lg n \rfloor} \frac{1}{2^i} \leq n(1 - 2^{-\lfloor \lg n \rfloor}) \leq n - 1$$

Theorem 3. After i rounds, there are $\lfloor \frac{n}{2^i} \rfloor$ trees with height i and size 2^i .

Proof by induction. Initially there are n trees of height 0 and size 1, so this is true for $i = 0$.

Assume the theorem is true for some i (induction hypothesis). Just before the $(i + 1)^{\text{th}}$ round, there are $\lfloor \frac{n}{2^i} \rfloor$ trees of height i and size 2^i . We can pair them up (if there are odd number of trees, leave the last one unpaired). When we union them, we get $\lfloor \frac{n}{2^{i+1}} \rfloor$ trees with height $i + 1$ and size 2^{i+1} (this doesn't depend on the value of **union.by**).

Therefore, the theorem is true by mathematical induction. \square

Theorem 4.

$$\left\lfloor \frac{n}{2^{\lfloor \lg n \rfloor}} \right\rfloor = 1$$

Therefore, after $\lfloor \lg n \rfloor$ rounds, there is one tree of height $\lfloor \lg n \rfloor$. Therefore, worst-case time complexity of **find** is $\Omega(\lg n)$.

2.5 Both union-by-rank and path-compression

2.5.1 Alt-Ackermann function

Definition 1. For $j \geq 0$ and $k \geq 0$,

$$A_k(j) = \begin{cases} j + 1 & k = 0 \\ A_{k-1}^{(j+1)}(j) & k \geq 1 \end{cases}$$

Here $A_k^{(0)}(j) = j$ and $A_k^{(i)}(j) = A_k(A_k^{(i-1)}(j))$.

Theorem 5. $A_k(0) = 1$

Theorem 6. $A_1(j) = 2j + 1$

Theorem 7. $A_2(j) = 2^{j+1}(j + 1) - 1$

Theorem 8. $A_3(1) = 2047$

Theorem 9. $A_k(j)$ is a non-decreasing function of k and j .

Theorem 10. $A_4(1)$ is way too large.

Proof.

$$\begin{aligned}
A_4(1) &= A_3(A_3(1)) \\
&= A_3(2047) \\
&= A_2^{(2048)}(2047) \\
&\geq A_2^{(2)}(2047) \\
&= A_2(A_2(2047)) \\
&= A_2(2^{2048} \times 2048 - 1) \\
&= 2^{(2^{2059}-1)} (2^{2059}) - 1 \\
&> 2^{2^{2059}} \\
&> 16^{16^{514}}
\end{aligned}$$

□

Definition 2. $\alpha(n) = \min(\{k : A_k(1) \geq n\})$

Theorem 11. $p < \alpha(n) \leq q \iff A_p(1) < n \leq A_q(1)$

2.5.2 level and iter

Let F be a DSU forest with n nodes. For a node x , let $x.p$ be its parent and $x.rank$ be its rank.

Theorem 12. $x \neq x.p \implies x.rank < x.p.rank$

Theorem 13. $x.rank \leq \lfloor \lg n \rfloor \leq n - 1$

We can partition the set of nodes into 3 parts:

- root nodes: $\{x : x = x.p\}$.
- leaf nodes: $\{x : x.rank = 0\}$.
- internal nodes: non-root and non-leaf nodes.

level and iter are functions which map an internal node x to an integer.

Definition 3. $\text{level}(x) = \max(\{k : A_k(x.\text{rank}) \leq x.p.\text{rank}\})$

Theorem 14. $k \leq \text{level}(x) \iff A_k(x.\text{rank}) \leq x.p.\text{rank}$

Theorem 15. $0 \leq \text{level}(x) < \alpha(\lfloor \lg n \rfloor + 1) \leq \alpha(n)$

Definition 4. $\text{iter}(x) = \max(\{i : A_{\text{level}(x)}^{(i)}(x.\text{rank}) \leq x.p.\text{rank}\})$

Theorem 16. $i \leq \text{iter}(x) \iff A_{\text{level}(x)}^{(i)}(x.\text{rank}) \leq x.p.\text{rank}$

Theorem 17. $1 \leq \text{iter}(x) \leq x.\text{rank}$

2.5.3 Potential function

Definition 5. For a node x , the potential function $\phi(x)$ is given by

$$\phi(x) = \begin{cases} \alpha(n) \cdot x.\text{rank} & x \text{ is a root or leaf node} \\ (\alpha(n) - \text{level}(x)) \cdot x.\text{rank} - \text{iter}(x) & \text{otherwise} \end{cases}$$

Theorem 18. $x \text{ is an internal node} \implies 0 \leq \phi(x) < \alpha(n) \cdot x.\text{rank}.$

To be continued ...